



Training-side signals per arm

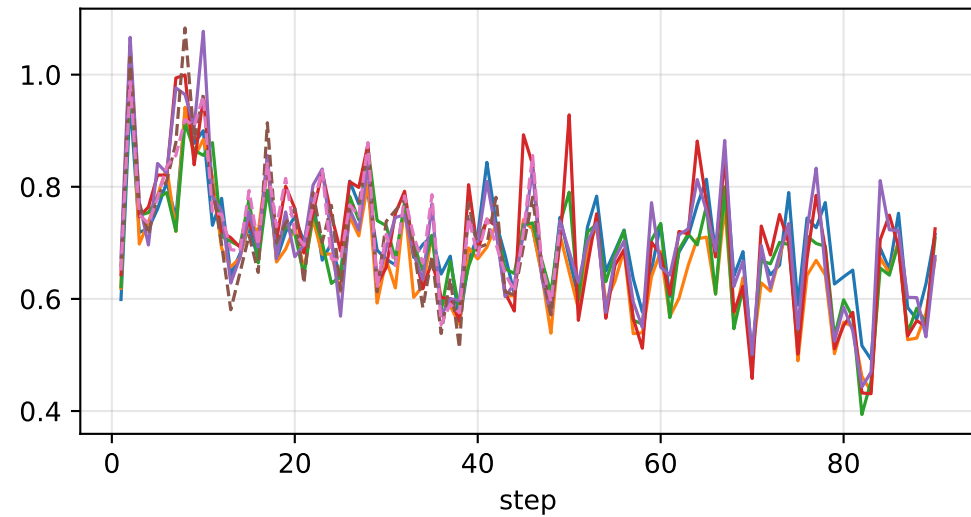
training reward (arm's own scale)



actor/kl_loss (target ~0.05)



policy entropy



response length (tokens)

